Algorithms Notes

Karl Pichotta

Spring 2013

Abstract

These are notes I've been taking for a graduate algorithms course taught by Greg Plaxton at UT Austin. They're taken primarily in-class, and probably only useful for my own personal reference.

1 Preamble: Useful identities

- For nonzero a, b, c, we have

$$a^{\log_b c} = c^{\log_b a} \tag{1}$$

- For large n, we have

$$\left(1 + \frac{1}{n}\right) \approx e \tag{2}$$

$$\left(1 - \frac{1}{n}\right) \approx e^{-1} \tag{3}$$

- For nonzero integer n, we have

$$n^2 = n + 2\binom{n}{2}. \tag{4}$$

  Both the algebraic and intuitive justifications for this identity should be pretty clear.

- For large n, we have

$$\sum_{i=1}^{n} \frac{1}{n} \approx \ln n + \gamma$$

  with γ the Euler-Mascheroni constant (about 0.58)

- For arbitrary positive n, k, we have

$$\binom{n}{k} \geq \left(\frac{n}{k}\right)^k$$

can be seen from the equation:

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$
$$= \frac{\prod_{i=0}^{k-1}(n-i)}{k!}$$

(FINISH ME)

2   1/22/2013: Chernoff Bounds, the Tails of the Binomial

If we have a binomaial RV $X \sim B(n,p)$, then we have the following Chernoff bound, for $0 \leq \delta \leq 1$:

$$P(X \leq (1-\delta)np) \leq \exp\left\{\frac{-\delta^2 np}{2}\right\}$$

Supposing we throw $100n \log n$ balls, by focusing our analysis on a single bin and then applying Union Bound, we get that the probability that some $X_i \leq \log n$ is $\leq n^{-44}$, using the above bound.

Fact: Thinking about Hashing, suppose we throw n balls into n bins, and we want a bound on the max load of any given bin. We might hope it's $O(1)$, but it isn't. Instead, it is

$$\Theta\left(\frac{\log n}{\log \log n}\right)$$

with high probability.

Upper bound: Define $Z = \max_i X_i$, with X_i the number of balls in bin i. Consider bin 1. $E[X_1] = 1$, but this tells us nothing about the tail of the distribution, which is our interest here.

What then, is

$$Pr\left\{X_1 \geq c\frac{\log n}{\log \log n}\right\}$$

Framing it in terms of Chernoff bounds, we will want to use:

$$\delta = c\frac{\log n}{\log \log n} - 1$$

We need to use the Large Deviations bound (bound (2) on handout). This is the following: Suppose $X \sim Binom(n,p)$, then

$$Pr\{X \geq (1+\delta)np\} \leq \left(\frac{e^\delta}{(1+\delta)^{1+\delta}}\right)^{np}$$

For concreteness, we take $c = 100$. Our $np = 1$, so the RHS of the above chernoff bound is

$$\ldots \leq \left(\frac{e^{\delta+1}}{\left(100 \frac{\log n}{\log \log n}\right)^{100 \log n/(\log \log n)}} \right) \leq \left(\frac{\log \log n}{\log n} \right)^{100 \log n/(\log \log n)} \approx n^{-100}$$

WE inflate the numerator by adding an extra e.

Sidenote: Now, this funny $\log \log n$ terms come from the fact that solving

$$x^x = n$$

gives you something very like $\log n/\log \log n$. This comes from canceling out lower order approximations to log.

Returning, by union bound, we argue that the probability that some bin gets $\geq 100 \log n/\log \log n$ balls is $\leq n^{-99}$.

Lower bound: So that was an upper bound. The lower bound is trickier. Let $k = \varepsilon \log n/\log \log n$, with ε a small positive argument. We won't be able to just reason about bin 1 and then use union bound—there is a constant probability $(1 - 1/n)^n \approx e^{-1}$ that bin 1 gets 0 bins. Let E_i denote the event that bin i receives at least k balls. We want to show that

$$Pr(\cup E_i) \geq 1 - \frac{1}{n^c}$$

This is, reminder, a lower bound on the max. That is, we're showing that

$$Z = \Omega(\log n/\log \log n).$$

We're showing that $Pr(E_1)$ is small, but it'll be useful.

So consider $Pr(E_1)$. We have

$$Pr(E_1) \geq Pr(X_1 = k)$$

(since RHS entails LHS) This is equal to

$$Pr(X_1 = k) = \binom{n}{k} (1/n)^k (1 - 1/m)^{n-k} \tag{5}$$

$$\geq \binom{n}{k} (1/n)^k (1/e) \tag{6}$$

$$\geq \left(\frac{n}{k}\right)^k (1/n)^k (1/e) \tag{7}$$

$$= \frac{1}{ek^k} \tag{8}$$

$$\approx \frac{1}{n^{\varepsilon'}} \tag{9}$$

We have the third to last step by the following useful lower bound:

$$\binom{n}{k} \geq (n/k)^k$$

which you see by expanding the formula and reasoning about the various values of quantities above and below the line.

Now, by linearity of expectation, we expect $n^{1-\varepsilon'}$ bins to get at least k balls.

What we want, however, is a high probability bound. That is, we want a statement of the form "at least one of bins, whp, will get at least k bounds". We can't use a Chernoff bound, because the E_i's are not independent, so therefore the distribution isn't Binomial.

From a high-level perspective: we throw n balls into n bins, then ask if E_1 occurred. If it did, then we'll be happy, because all we want is one of the E_is to occur (NB $PR(E_1) \geq 1/(n^{\varepsilon'})$). If E_1 doesn't occur, then note

$$PR(E_2|E_1^c) \geq Pr(E_2) \geq 1/n^{\varepsilon'}$$

and so on with all of the $Pr(E_i)$'s. We can therefore argue that the probbability that noE_i occurs is

$$\leq (1 - 1/n^{\varepsilon'})^n$$

We want to factor this last equation into

$$... = ((1 - 1/n^{\varepsilon'})^{n^{\varepsilon'}})^{n^{1-\varepsilon'}} \tag{10}$$

$$\approx (1/e)n^{1-\varepsilon'} \tag{11}$$

Instead of having $1/e$ to some logarithmic quantity, we have it to some polynomial quantity. So this is much, much less than inverse polynomial bound, and we've more than satisfied our sharp threshold.

An exercise close to the one we did at the begnning—using Chernoff bounds, we can argue that the number of flips of a fair coin to get $\log n$ heads with high probability is

$$\Theta(\log n)$$

(we use Chernoff bound (1), as in the first example). This comes up in at least one way to analyze randomized Quicksort.

In last class, we saw that the expected number of comparisons is $\Theta(n \log n)$. We now argue that, with high probability, the number of comparisons is $O(n \log n)$. Note this is quite different: what we say now is that the probability that you exceed the expected runtime by a factor of, say, 50, is very small. Note these RVs aren't binomial, but we'll be able to bound their behavior with binomial RVs.

First, we'd LIKE to show the number of comparisons involving a specific key is $O(\log n)$ with high probability. This is not true—there is a $1/n$ chance of an element being the first pivot, and it gets compared to n elements.

So what we'll do instead is to use a charging scheme: we charge a comparison to a non-pivot. Whenever we make a comparison, "charge" the comparison to the nonpivot (each comparison involves a pivot and a nonpivot). It IS true that the charge to any key is $O(\log n)$ with high probability. Once we show that, it immediately follows that the total number of comparisons is $O(n \log n)$ WHP. HOw can we use the prvious result to convince ourselves of this?

4

3 1/24/2013: Hashing

Let's think about Randomized quicksort. We know that randomized quicksort has expected $O(n \log n)$ behavior. (that is, there are no "bad inputs" for it in expectation). Note also that Randomized quicksort has $O(n \log n)$ runtime with high probability (that is, we'll give a nice bound on the probability of the runtime being asymptotically higher).

We can think about quicksort from the perspective of a particular fixed key, call it x. Recall the charging scheme: we "charge" each comparison to the nonpivot (which we can do, since each comparison is between a pivot and another value).

Claim: The key x gets $O(\log n)$ charge with high probability (the failure probability will be $1/n^c$ for arbitrarily large c. (If we prove this, then by Union Bound, we have the total charge to all n keys is $O(n \log n)$ WHP.)

Proof of Claim: We have a "herd" of keys, initially of size n, and select one to be the pivot. The pivot is remarkably lucky—it gets 0 charge. If x is chosen as the pivot, that's nice. If, on the other hand, x is not chosen as pivot, then it will get charge 1. Then the keys will be partitioned into two "herds", one of which will never again be compared with x.

Consider the process where we start with a positive integer n. We flip a coin. If we get heads, we set the number to an integer uniformly drawn from $[0, 3n/4]$ (floored). If we get tails, then we set the value uniformly to between 0 and $n - 1$. So at each stage, the value gets smaller (by at least one). The process stops when the number gets to 0.

We use this process as follows. Picking a pivot at random, there is a 0.5 probability that the pivot sits in between $n/4$ and $3n/4$ in the sorted list (it sits in the middle half). There's also a .5 chance it's not in the middle. The number in the process is just the size of the "herd" x is in.

If we get a good pivot, then the herd value is at most $3n/4$ (imagine getting the rightmost pivot in the middle half, say). This is like flipping a head. In the worst case (flipping a tail), then the size of the herd decreases by 1.

The process can't involve more than some constant times $\log n$ heads: each head decreases the herd size by at least $3/4$. (Prove this?) So x sits there and hopes for good pivot choices from its perspective. (This is basically the end of the proof.)

We don't have a binomial variable, but we can relate it to this simpler process defined above, and use a Chernoff bound on that.

3.1 Hashing with chained overflows

We have a hashtable, visualized as an array of buckets, numbered $0, \ldots, k - 1$. We have a has function
$$h(x) \in \{0, \ldots, k - 1\}.$$
In order to handle collisions, we can use a linked list to chain them together in the bucket. Searching for an item in the table later involves computing the hash and then traversing the linked list.

Suppose you put k things in. You're generally hoping that each bucket tends to have $O(1)$ elements. In the worst case, you traverse a linked list (and have $O(k)$ lookup). In the best case, we get $O(1)$ lookup.

What's the runtime of h? Supposing we have n keys, each of which is $10 \lg n$ bits. (Note we're using the RAM model: the word size of the machine we're programming is logarithmic in the input size in bits of the problem. So if we have a million-bit instance, we assume that we can manipulate words of log 1m in constant time.) So that means that our word size is $\Theta(\log n)$ bits. That is, these keys will fit into a constant number of words. We'll be relying on this later to assume various operations are $O(1)$.

So for now, we think of h as taking any key and mapping uniformly random from $0, k-1$. This is of course not quite right: h must be deterministic. However, we think of the hash for our purposes as choosing uniformly at random.

This relates directly to the bin/bucket problems we discussed previously.

Claim: The average time for a search is $O(1)$, assuming n buckets (with n keys).

No proof (prove this?). Basically, the vast majority of buckets will get very few elements.

On the other hand, what's the expected max search time? That is, has n keys into n buckets with our idealized hash. The worst-case search time is the longest linked-list we get. We looked at this last time: The max load is

$$\Theta\left(\frac{\log n}{\log \log n}\right) \ w.h.p.$$

So it's only marginally better than using a red-black tree.

3.2 Perfect Hashing

At a high-level, we use *Perfect Hashing* to obviate this non-constant load problem. That is the following:

We assume we're dealing with a *static* set of n keys. That is, we construct a hash-table structure that is specific for the particular n keys that we have at hand. As before, assume each key is $10 \log_2 n$ bits.

Desiderata:

- We want to construct a hash table with $O(1)$ worst case search time. NOTE this isn't just a matter of expectation: we want guarantees about the worst case.

- We also want to use $O(n)$ space (that is, our solution can't be "use $O(2^n)$ keys").

- Further, we want "fast" construction

A naive approach is the following: Repeatedly pick a new hash function until you find one inducing max load of $O(1)$. We can do this because we have static keys (we know them in advance). Eventually we'll find one with $O(1)$ worst

case. However, this won't give us fast construction: we'll have to run this an exponential number of times in order to find this (I don't quite get the proof, but it involves one of the results we showed last time involving $\log n$).

We'll use a twist on this: we'll follow a similar approach, but our criterion for selecting a hash function will be looser. we'll use a two-tiered approach: What we'll do is count "collisions", and as long as we don't have more than n collisions, we'll call it good. So we have n buckets, but instead of a linked list in each bucket, we have a hash table at each of the n buckets. That is, we have one primary hash table, and n secondary hash tables.

We want the total size of all these tables to be $O(n)$. We do our primary hashing; based on how many elements hash into a location, each bucket has a hash table of that size. We'll have the size of a secondary hash table being quadratic in the number of elements there. So if we have 20 elements in a bucket, we'll have a hash table of size approximately 20^2. Why? We want to totally avoid collisions at the second level of hashing. The thing about quadratic size is that this is the threshold at which we'll suffer 0 collisions in the secondary hash tables (we can repick secondary hashes).

One concern is that since we're wasting space in the secondary tables, when we add the size of the secondary tables up, it'll be too large. So we need to show that the sum of the table sizes will be linear (intuitively, we'll show the vast majority of buckets get only a constant number of keys).

So OK getting to details. What is a "collision"? Mapping keys to buckets, we'll let Y_i denote the RV corresponding to the number of keys mapping to bucket i at the top level (for $1 \leq i \leq n$). Let X be the RV denoting the number of collisions. Then

$$X = \sum_{i=1}^{n} \binom{Y_i}{2}.$$

That is, if 10 keys are mapped to a bucket, that gives 10 choose 2 pairwise collisions. Note that

$$X = \Theta\left(\sum_{i=1}^{n} Y_i^2\right)$$

which is, magically, the total size of the secondary hash tables. So X is constant-factor-related to the space required for the secondary hash tables. (Note that it's good enough, then, to pick a hash function that induces at most $100n$ or $1000n$ collisions.)

Now, to get at $E[X]$, we express X as a sum of indicator variables. Define Z_{ij} to be the indicator variable taking 1 if the i, jth keys collide, and 0 otherwise.

Note there are n choose 2 such vars. So

$$E[X] = \sum_{i,j} E[Z_{ij}] \tag{12}$$

$$= \sum_{i,j} 1/n \tag{13}$$

$$= \binom{n}{2}/n \tag{14}$$

$$= \frac{n-1}{2} \tag{15}$$

where we appeal to the fact that we have an idealized hash function. This is less than $1/2$ of our target of getting $\leq n$ collisions. (It is important that we choose a constant $> 1/2$ when defining our criterion for accepting a hash functions.)

Call a hash function "good" if it induces $\leq n$ collisions; "bad" otherwise. On average, the number of collisions is $n/2$. We can therefore bound the percentage of all hash functions that are "bad". If, for example, 90% of the has functions are bad, then the expected number of collisions would be at least $0.9n$. So we have that at most half of the hash functions are bad. Note that this is a huge overestimate: the only way this can happen (that is, we get the expected value we had before) is if all the bad hash functions have exactly n collisions and the good ones have 0.

So in our first phrase, picking a primary hash function, This is, in the worst case, like flipping a fair coin until you get a heads. So the expected number of trials is $O(1)$ (it is Geometric(0.5)). So, with high probability, the number of trials is $O(\log n)$ (proof?).

Once we have the top-level hash function, then, the Y_i values are determined. In bucket i, we use a hash table of size Y_i^2. We repeatedly pick hash functions for bucket i until we find one that gives no collisions at all.

That is, the secondary hash table at buekct i has Y_i^2 size, Y_i keys. The expected number of collisions is calculated using a similar approach before. We define $\binom{Y_i}{2}$ indicator variables, and the probability that one of those is 1 is $1/Y_i^2$. So defining C_i th enum of collisions at i, we have

$$E[C_i] = \binom{Y_i}{2}(1/Y_i^2) = \frac{Y_i - 1}{2Y_i} \leq \frac{1}{2}.$$

The analysis is identical for every bucket. So cool! We're done, this scheme works.

3.3 Realistic Hash Functions

The tricky thing here is that we've been assuming that we have these idealized hash functions that distribute uniformly, and we can generate them easily (and they're independent). What is it like in real life?

When we were talking about idealized hash functions, the associated family corresponds to the family of

$$n^{n^{10}}$$

hash functions (this is the number of functions from $10 \lg n$ strings to n things, I think?). Note if we represent these naively, then we use base-2 log num bits to specify an element. This is probematic, because this description of a function is of length $n^{10} \lg 10$ bits.

The key thing to note here (that'll get us around this) is that pairwise independence is sufficient for the family of hash functions we use. What do we mean? We mean the following: A family of hash functions $\mathcal{H}$ is pairwise independent if, for any distinct keys x, y, and any (possibly equal) bucket indices i, j, if the hash function h is drawn uniformly at random from $\mathcal{H}$, then

$$Pr(h(x) = i, h(y) = j) = \frac{1}{B}$$

with B the number of buckets.

This is sufficient because, in our analysis, we were only interested in pairwise indicator variables. We were never interested in any more complicated conditions.

We're interested in designing a family of functions mapping from $(10 \log_2 n)$-bit strings to $(\log_2 n)$-bit strings.

We saw that for analysis, we didn't need full independence, but just pairwise independence. This is because our analysis cared only about the number of pairwise collisions, which we can write as a sum of indicator variables X_{ij}, which is 1 iff i and j collide. Recall that $P\{X_{ij} = 1\} = 1/p$, with p the number of buckets.

In our analysis, in fact, it's fine to have approximate pairwise independence; that can give us OK bounds. So we want a function h such that, for two distinct keys $x \neq y$ (and any i, j, not necessarily distinct), then

$$Pr\{h(x) = i \wedge h(y) = j\} = O\left(\frac{1}{n^2}\right).$$

IN the above, i and j are bins; x and y are keys.

First, we'll pick a prime p a bit bigger than n^{10} (why?). The prime number theorem tells us that about, picking things around the neighborhood, we only have to try a logarithmic number of keys (in n^{10}) in expectation before finding a prime. (There is also a theorem indicating that there is definitely a prime between k and $2k$; in particular, this is n^{10} and $2n^{10}$).

So let's hash from $\mathbb{Z}_p$ to $\mathbb{Z}_n$. Consider

$$h(x) = [ax + b \mod p] \mod n$$

with a and b chosen from $\mathbb{Z}_p$. (Note that we have a and b because we want it to be the case that, under our analysis, there will be no "bad input"—if we have an adversary picking keys, they won't be able to pick bad keys if they know

what the family looks like; if there were no a and b, then they would be able to do so. In other words, our family needs to have more than one hash function in it if we want to argue a small number of expected collisions.) This isn't quite uniform between 0 and $n - 1$, but we're quite close. The inner hash

$$ax + b \mod p$$

is, I think, actually uniform over $\mathbb{Z}_p$ (proof?). The outer hash, then, will be approximately uniform over $\mathbb{Z}_m$ (proof?). This satisfies the condition for approximate pairwise independence that we defined earlier.

So at the top level, we pick a p, then we repeatedly choose a and b; for each a and b, we check how the particular hash function works. We repeatedly do that until we get a good a and b; we store them, as they fully parametrize a hash function. We do this for each primary hash function and each secondary hash function.

4   1/29/2013: Dynamic Programming

4.1   Examples

4.1.1   contiguous subarrays of booleans

Suppose we're given an $n \times n$ boolean array A. We want to find the side-length of a largest all-true contiguous square subarray of A. So if there's a 3×3 subarray of Trues, then the desired answer is at least 3.

Now, certainly this problem is solvable in polynomial time. For each side length k, there are only n^2 different places where the $k \times k$ subarrays top-left corner could be; we could just exhaustively check this for all k and compute the answer.

We can get a better polynomial bound using Dynamic Programming. What DP does is, instead of solving a single problem instance, find a bunch of problem instances to solve, and do so in such a way that the larger problem instances can use the computations performed for the smaller instances.

So let a_{ij} be the size of the largest all-true contiguous square subarray with lower right corner at location (i, j). This is a family of n^2 subproblems. Our answer to the original problem is just

$$\max_{i,j} a_{ij}.$$

Now, what is a good order to solve these n^2 subproblems? What we'd like is that, whenever we get to a particular problem in this ordering, we can very easily solve it in terms of the instances we've already solved. In other words, we want to be able to write a suitable reccurrence for the a_{ij}'s. Thinking about it, it's not too tough to see the following works:

$$a_{ij} = \begin{cases} 0 & \text{if } (i,j) \text{ entry is F} \\ 1 + \min(a_{i-1,j}, a_{i,j-1}, a_{i-1,j-1}) & \text{otherwise} \end{cases} \tag{16}$$

So filling in left-to-right row by row (or top-down column by column) should work. The important thing is that, when we get to a position, the spot above it, to the left of it, above and to the left of it, are filled in. (A small note is that we need to interpret out-of-bounds indices as 0 above.)

So the above algorithm runs in $O(n^2)$ time, which is a nice, much-faster polynomial-time algorithm than the brute-force poly-time algorithm.

4.1.2   Rod-cutting problem

Commonly, Dynamic programming yields polynomial-time algorithms for problems where the brute-force technique is exponential. For example, suppose we have a rod of length n inches ($n \in \mathbb{Z}^+$). We can cut it into any integer-length pieces (such that they sum to n). For every i, we have a price p_i that we can sell a rod of length i for. The problem, then, is to cut up the rod into pieces in such a way that we maximize the total prices for the pieces.

When we cut the number up, we get a multiset of positive integers; in the language of number theory, this is a *partition* of n. If the number of partitions of n were sufficiently low, we could just generate them and calculate their value; however, the number of partitions grows superpolynomially. Consider the numbers
$$1, 2, \ldots, \lceil 2\sqrt{n} \rceil - 1, \lceil 2\sqrt{n} \rceil.$$
Pair up the first two numbers and the last two numbers:
$$1, 2, \ldots, \lceil 2\sqrt{n} \rceil - 1, \lceil 2\sqrt{n} \rceil.$$
For the first $\lceil 2\sqrt{n} \rceil + 1$ elements, you can either partition them into the two outermost pieces or the next innermost pieces.

We repeat for $3, 4$ and the next two inner numbers on the right. For each of these, we eat up $\sqrt{n}$ of the rod. We'll have $\sqrt{n}/2$ such decisions for what to do with the pieces of length $\sqrt{n}$, and we'll have about
$$2^{\Omega(\sqrt{n})}$$
ways to partition the rods up. (In fact, it ends up being $2^{O(\sqrt{n})}$ too.) So the number of partitions of n is pretty clearly superexponential.

We can solve this, however, with a simple one-dimensional dynamic program. For each i, let v_i be the maximum value of a rod of length i. We will use the following recurrence for v_i in terms of the smaller v_j's. For a partition of i, there will be some length ℓ of the last partition. So
$$v_0 = 0$$
$$v_i = \max_{1 \le \ell \le i} v_{i-\ell} + p_\ell$$

We get $O(n^2)$ time of this: we have n subproblems, but the ith subproblem takes $O(i)$ time to solve (summing over these subproblems uses the familiar $\sum_i i^2 = O(n^2)$ identity).

4.1.3 Longest Common Subsequence

We have two sequences X, Y of symbols over some finite alphabet. Writing X as

$$X = x_1 x_2 \ldots x_n,$$

we define a subsequence of X as a subset of entries x_i, concatenated in order. We want to find the largest integer k such that there are subsequences in X and Y of length k such that both subsequences are the same.

The brute-force approach here is exponential: it's pretty easy to see (there are 2^n subsets of n sets).

We use the familiar two-dimensional DP solution. Let X_i denote the length-i prefix of X:

$$X_i = x_1 \ldots x_i,$$

and similarly with Y_j. Now, let a_{ij} be the length of the LCS of X_i and Y_j.

Supposing $|X| = m$ and $|Y| = n$, we have $m \times n$ subproblems. The final answer to the question is a_{mn}. So we have the base cases

$$a_{0j} = 0$$
$$a_{i,0} = 0$$

for all i, j. For the recursive case, we have

$$a_{ij} = \begin{cases} a_{i-1,j-1} + 1 & \text{if } x_i = y_j \\ \max\{a_{i,j-1}, a_{i-1,j}\} & \text{if } x_i \neq y_j. \end{cases} \tag{17}$$

4.1.4 The Partition Problem (a pseudopolynomial algorithm)

This problem is NP-hard. Suppose we have n positive integers $x_1, \ldots, x_n$. We want to know whether the x_i's can be partitioned into two multisets of equal sum.

We give a DP algorithm. First, let $\sum x_i = 2s$. (Note if the sum is odd, we return that there is no such partition). To determine s, it' a natural thing to consider computing

$$a_{i,j} = \begin{cases} T & \text{if there's a subseq. of } x_1, \ldots, x_i \text{ summing to exactly } j \\ F & o.w. \end{cases} \tag{18}$$

Note $\forall i, a_{i,0} = T$. Further, $\forall j > 0 \; a_{0,j} = F$. When $i, j > 0$, we have

$$a_{i,j} = \begin{cases} a_{i-1,j} & \text{if } j < x_i \\ a_{i-1,j} \lor a_{i-1,j-x_i} & o.w. \end{cases} \tag{19}$$

NB we only really need the first term for the technicality of the subscript on the RHS disjunct being negative.

Now, the number of table entries is $O(nS)$. Is this polynomial time? Intuitively, no—S can be very large. We argue that the algorithm is not polynomial;

to do so, we need to find one family of inputs where the runtime is not upper-bounded by a polynomial function of the input in bits.

What do we mean by polynomial time? Well, we upper-bound the runtime of an algorithm according to a polynomial function of its input in bits. What's the input of this like? Well, consider an input where each x_i is an n-bit integer. The input size for this problem is $\Theta(n^2)$ bits (we have n n-bit numbers). For us to claim that we have a polynomial runtime, we need that, for such inputs, the runtime is polynomial in n. S, however, is $O(n2^n)$, since an n-bit integer can be as large as 2^n in value (and we have n of them). So, more importantly, S can be $\Omega(n2^n)$. So this algorithm, which kinda looks like it's polynomial, actually isn't polynomial in the input size.

However, sometimes these sorts of algorithms are useful. If we know, for example, that each of the numbers is at most n^3, then we have that S is at most n^4, and we have a polytime algorithm. We call algorithms like this *pseudopolynomial*: they are polynomial in the input size if the input integers are represented in unary.

5   1/31/2013: Greedy Algorithms

5.1   A Scheduling Problem

Suppose we have n tasks, each with a positive integer deadline and an execution requirement. So an input may look like

task	Deadline	execution requirement
A	8	5
B	6	2
C	9	2

We want to know if we can meet all the deadlines in a nonpreemptive schedule? So, in the above, suppose we run B at time 0 and C at time 2; both of these take time 2, and have met their deadlines. If we run A next, though, it'll miss its deadline. However, if we run in the order BAC, we get a feasible schedule.

We may want to try to get at this by Dynamic Programming, but we don't need it: a greedy solution ends up sufficing. We use *earliest deadline*, breaking ties arbitrarily. This will yield a feasible schedule if one exists.

We want to argue that the answer to the greedy algorithm is right iff there's a feasible schedule. Clearly, Left to Right is easy. The argument is R to L, then: If there's a feasible schedule, then we need to show the greedy algorithm finds it.

So take a feasible schedule. Assume the feasible schedule has no "spaces": there's no advantage to idle time, so everything is "squeezed" as far as possible to the left. Suppose that we have a schedule $ABCD$ of four elements, in that order. Suppose that C has an earlier deadline than B; earliest deadline would have chosen the order $ACBD$, rather than $ABCD$.

We can modify the schedule by leaving everything the same except for B and

C, swapping the latter two. Note the new schedule is clearly the same length. We want to show that $ACBD$ is also feasible. Clearly C meets it deadline still— it's moved earlier. We need to show that B is meeting its deadline. Since, by assumption, C had a more stringent deadline than B, it must be the case that if B finishes when C did before, it certainly meets its deadline, since it meets C's deadline. In other words, before, with $ABCD$, C was meeting its deadline; we know that B's deadline is later than C's, so in $ACBD$, B definitely meets its deadline.

We apply the same logic to the new problem instance to argue that it's feasible (eventually we will get to the earliest-first schedule).

5.2 A variation of the scheduling problem

Suppose that, in addition to the deadline and execution requirements, each task has a positive profit $p_i > 0$ associated with it. Our goal is to find a max-profit feasible subset (that is, we won't in general use the entire set of tasks). What we'll do is combine the greedy approach with dynamic programming.

We first order tasks by deadline (in nondecreasing order). We then consider all prefixes with respect to that particular ordering. The DP's subproblem is $a_{k,t}$, defined as the max profit we can obtain with a schedule using only the first k tasks and time $\leq t$. We can write a recurrence fairly straightforwardly (left as an exercise).

The runtime of the above will be polynomial if the execution requirements are polynomially-bounded. That is, if every task has a task that is $O(n^c)$ for some c. Without such an assumption, we could get an exponentially-sized table in our DP

5.3 Matroids

Captures a large class of greedy algorithms. Consider the MInimal Spanning Tree problem, in particular, Kruskal's greedy algorithm. One way to convince ourselves that Kruskal's algorithm is correct is to reason directly about minimal spanning trees. Another is to frame the algorithm as a matroid problem, and use a general property about Matroids.

A Matroid is a pair

$$M = (S, \mathcal{I})$$

with S analogous to the vertices in a graph: it's just a set. $\mathcal{I}$ is called the *independent sets* of S: $\mathcal{I} \subseteq 2^S$ (elements are subsets of S). We need two more properties:

1. **Hereditary Property**: if $A \in \mathcal{I}$ and $B \subseteq A$, then $B \in \mathcal{I}$.

2. **Exchange Property**: If $A, B \in \mathcal{I}$ and $|A| > |B|$, then there is some $x \in A \setminus B$ such that $B \cup \{x\} \in \mathcal{I}$.

The hereditary property tells us that removing elements from an independent set yields an independent set. In particular, this tells us that $\emptyset \in \mathcal{I}$.

Define a maximal independent set as an independent set such that, if you add any additional elements to it, it will no longer be independent. The exchange property tells us that all maximal independent sets will have the same cardinality. A maximal independent set is sometimes called a **basis**

In a **weighted matroid**, each element $x \in S$ has an associated weight $w(x)$. In applications, we often want to compute a maximum (or minimum) weight maximal independent set. For example, eventually we'll have a minimal spanning tree representing a minimum weight maximal independent set (the maximal independent sets will correspond to the spanning trees).

5.4 Matroid Greedy Algorithm

Suppose we want to do maximization.

- First, sort the elements of S in nonincreasing order of weight (if we are minimizing, we'll sort in the other direction).

- Initialize $A := \emptyset$

- For each $x \in S$, in the order determined in the first step:

    - If $A \cup \{x\} \in \mathcal{I}$, then $A := A \cup \{x\}$.

And that's it. Thinking about Kruskal's algorithm, the set $\mathcal{I}$ will be the set of acyclic graphs—the conditional above corresponds to the conditional in that algorithm.

How do we show this algorithm is correct? First, index the elements of S $1, \ldots, n$ according to the ordering in the first step. Call the weight for element i in this ordering w_i. For each i, calculate some max-weight maximial independent set B. Mark down if element i is in B or not. Also, for each i, consider A, the set computed at that step by the algorithm. What we need to do is prove, for each step, A is a maximal weight independent step. Similarly, look at whether elem i is put in A.

Look at the first part where the sets A and B differ. It can't be the case that the elem i is in B but not in A (proof? It has somethign to do with the hereditary property, not sure what). So the first place where the two differ is that an element i must be in A and not in B. WE note the first place i where A and B differ. Construct $A' \subset A$ from all the before i in the ordering, and also elem i. Now, we know $|B| \geq A'$, since $|A'| \leq |A|$, and $|B| = |A|$, by the exchange property.

We repeatedly apply the exchange property to add elements from B to A'. The exchange property tells us that we can add an element from B to A' and get an independent set, so long as $|A'| < |B|$. This process terminates exactly when $|A'| = |B|$.

When we're done growing A', B gave A' all its elements from $i + 1$ onwards except for one of them. In other words,

$$A' = (B + i) = j$$

for some j. That is A' has elem i while B doesn't, and A' deosn't get some $j > i$. However, since the elements are ordered by weight, and $w_j \leq w_i$, we can conclude that
$$w(A') \geq w(B).$$

Now we play the whole game over, but with B replaced by A'. This gets us "one step" in the right direction, and we iterate this process until we end up with $A' = A$, at which point we can argue that A was a max-weight maximal independent set.

5.5 Some examples of Matroids

- **Uniform matroid**: Let $|S| = n$. Fix $0 \leq k \leq n$. $\mathcal{I}$ is the set of all subsets of S of size $\leq k$. This is a matroid. The hereditary property is pretty obvious. The exchange property is also pretty clear: if we have A and B such that $|A| < |B|$, then there's clearly an element in B that's not in A such that adding it to A will result in an element of cardinality $\leq k$.

- **Partition matroid**: Fix a partition of S into $\{S_1, \ldots, S_k\}$. Define $\mathcal{I}$ to be the set of all subsets $A \subseteq S$ such that $\forall i |S_i \cap A| \leq 1$.

  A basis (maximal independent set) for this matroid would be a set that chooses exactly one element from each S_i. The number of such bases is

$$\prod_{i=1}^{k} |S_i|.$$

  The hereditary property is pretty clear. The exchange property is a bit trickier, but also pretty clear.

- **Scheduling problem**: we have n tasks numbered from 1 to n. Task i has a positive integer deadline d_i, a positive weight w_i (the profit of the task, say), and an execution requirement of 1 (unit tasks). We're scheduling these tasks on a shared resource.

  What is the maximum-weight schedulable subset of the tasks (in general, we might not be able to meet all tasks of the schedule)? This is a bit complex: Suppose there are 2 tasks with deadline 1, and two tasks with deadline 2. We can schedule at most one of the former, but two of the latter (but only if we schedule neither of the first two).

  This problem ends up having a useful matroid structure, allowing us to use the matroid greedy algorithm.

  Define the set S as the set of n tasks. $\mathcal{I}$ will be the schedulable subsets of the tasks. If this works, our optimal solution will certainly be a maximal independent set (this isn't necessarily true if we have negative weights I don't think).

  Now, to show $(S, \mathcal{I})$ is a matroid. The hereditary property is pretty clear. The exchange property is a bit trickier. Suppose $A, B \in \mathcal{I}$, with $|A| > |B|$.

We proved earlier that A is schedule iff any ED schedule (???) of A is feasible (wtf does this mean?). Define a "canonical" ED schedule as one of the ED schedules: break ties in a particular manner. That is, we're constructing a schedule of a task A; we list tasks in increasing order of their index or something like that, guaranteeing a unique tie-breaking schedule. The schedule, then, will be some ordering

$$A = x_{i_1}, x_{i_2}, \ldots, x_{i_{|A|}}$$

and

$$B = x_{i_1}, x_{i_2}, \ldots x_{i_{|B|}}.$$

Note that there will be no "empty spaces" in the schedule.

Let's focus on the last task appearing in A that isn't a part of B. Define z to be this tasks: the last task in the canonical ED schedule of A that does not belong to B. To show the exchange property, we need to show there's something in A that can be added to B. We'll claim that z is such an element; that is, $B + z \in \mathcal{I}$ (that is, $B + z$ is schedulable).

- **Case 1**: z is the last task in the ED schedule of A. Because B is shorter, we can just tack z onto the end of B; it will surely meet its deadline, because it will start at least as early as it started in A. This isn't necessarily an ED schedule, but is feasible.

- **Case 2**: The ED schedule of A ends with $z, y_1, \ldots, y_\ell$ with $\ell \geq 1$ (that is, z isn't the last element of A's ED schedule). What we do is jimmy around with the schedule of B to arrive at a feasible schedule for $B + z$. Leave the schedule before y_1 alone; put z in the spot where y_1 is. Find y_2 in B's schedule, and put y_1 in it; similarly, put $y_{\ell-1}$ into the schedule where y_ℓ was, and tack y_ℓ onto the end (recall, these are the occurrences of $y_1, \ldots, y_\ell$ in B's ED schedule. Now, y_ℓ in the new schedule for B is no further to the right than it is in A's schedule. Similarly, the furthest to the right that $y_{\ell-1}$ could possibly be is where it is in A's schedule. So this is a feasible schedule (KBP: i don't see why $y_1, \ldots, y_\ell$ are in B's schedule. Maybe they don't have to be?).

- **Scheduling with release times**: a more complicated version of the above. Assume each task also has a positive integer **release time**. Previously, we assumed all tasks were available at time 0; we restrict that assumption now: tasks become available to executable at a certain point in time, and can only be scheduled past that time.

  Even with this additional constraint, we can construct a matroid structure (we don't here).

- **Transversal matroids**: even more general. Say we have a bipartite graph (U, V, E) (with U, V the two sets of vertices such that $\forall (u, v) \in E$, u and b aren't both in U or V). Consider the matroid $(U, \mathcal{I})$, with every

$S \in \mathcal{I}$ a "matchable" subset of U, defined below. A **matching** is a subsets of the edges E that induces degree at most 1 on every vertex. That is, in a matching has no vertex incident on more than 1 edge (that is, every vertex has degree ≤ 1). Given a matching, call a vertex **matched** if it's incident on one of the edges in the matching. A subset U of vertices is **matchable** if there exists some matching matching every $u \in U$.

So our matroid is $(U, \mathcal{I})$, with U the same set in our bipartite graph $(U, V, \mathcal{I})$ (proof that this is a matroid?). Relating to the last scheduling problem: Consider U the set of tasks, and V the set of times during which a task can start. So supposing task i has deadline of 5 and a release time of 2 (and takes 2 time units), we'll add edges from $i \in U$ to $2, 3, 4$ in V. Given this setup, every matchable subset corresponds to a valid scheduling.

The hereditary property is easy: if a set of tasks is matchable, of course any subset is matchable.

The exchange property is tricky: say we have $A, B \in \mathcal{I}$, $|A| > |B|$. Let M_A be a matching matching exactly A (which exists by assumption), and let M_B be a matching for B. That is, the set of nodes matching in M_A is exactly A (so $|M_A| = |A|$), ditto B. We need to show that there is some M matching $B + x$, for some $x \in A \setminus B$. Think of these matchings as sets of edges.

Consider the graph $G \equiv (U, V, M_A \Delta M_B)$ (taht is, the symmetric difference of M_A and M_B: edges in exactly one of M_A, M_B). Because M_A is such that each node is of degree at most 1, we have that $M_A \Delta M_B$ induces degree of at most 2. G consists of vertex-disjoint paths and cycles (the cycles are of even length; each cycle would alternate between an edge in M_A and and edge in M_B—you can't have two consecutive edges from M_A, because this would give us a node of degree 2).

We cannot have only cycles: since we know that $|M_A| > |M_B|$, we must have some paths (since every cycle draws an equal number of edges from M_A and M_B, I think). In particular, there must be at least one odd-length path that begins and ends with an edge in M_A. Note that the paths, like the edges, alternate from edges in M_A and edges in M_B. So assume WLOG that this path starts in U and ends in V (it's odd-length, and a bipartite graph, so an odd-length path can't begin and end in U or V). Call the node in U that the path starts at x.

We can obtain a matching for $B + x$ by making a slight hack on M_B. Take the odd-length path above; get rid of the original edges from M_B, and add the edges from M_A that were touching the same nodes. This doesn't lose matchability (every one of the nodes on the path is still connected to exactly one edge, this time the edge from M_A), and also now x is matched.

It's worth noting that the computational problem to solve the matroid greedy algorithm is trickier for this formulation: we have to determine matchability for arbitrary subsets, which is a little tricky.

- **Graphic matroids**: Fix an undirected graph $G = (V, E)$. Define $M = (S, \mathcal{I})$, with $S = E$ and $\mathcal{I}$ the set of all acylic subsets of E. The hereditary property is pretty obvious: removing edges won't add cycles.

We don't have time for the exchange property in class.

6   2/7/2013: Minimal Spanning Trees and Matroids

Recall that, given an undirected graph $G = (V, E)$, we construct a graphic matroid from it by the set $(E, \mathcal{I})$, with $A \in \mathcal{I}$ iff A is acyclic. Here, the Maximally independent set are exactly the spanning forests; if G is connected, then it is a spanning tree.

Note that a **minimal dependent set** of a matroid $(S, \mathcal{I})$ is a dependent set (that is, a subset of S not in $\mathcal{I}$) such that if we remove any element from it, we get an independent set. If we have a graphic matroid, then a minimal dependent set will be a cycle (we can't throw any edge out of the set without getting something that's acyclic).

6.1   A theorem and some corollaries about Matroids

Lemma: let $M = (S, \mathcal{I})$ be a weighted matroid. Let A, B be distinct max-weight bases of M. Then there exists a sequence $A_1, \ldots, A_k$ of max-weight bases of M such that $A = A_1, B = A_k$, and for every i with $1 \leq i < k$, we have

$$|A_{i+1} \setminus A_i| = 1.$$

That is, A_{i+1} and A_i differ by a "swap":

$$A_{i+1} = A_i + x - y$$

with x and y having equal weight (since they're all max-weight).

   Proof: How do we prove this? One way is to use something similar to our proof of correctness for the matroid greedy algorithm. We order the elements (x_i, w_i) by decreasing order of weight, and then we line up A and B, marking whether $x_i \in A$ and $x_i \in B$. We look at the first element such that $x_i \in A$ but $x_i \notin B$, and, we obtain B' by growing an A^* somehow (how?!), starting with A^*, and growing to a basis via repeated application of the exchange property with B. So we'll end up with B', which will differ from B by a single swap. There'll be one element that B does have but that B' doesn't have; other than that, B' (which is, I think, what you get after constructing as many A^*s as you can) and B will be identical.

   Since A and B are both max-weight bases, what can we conclude about the basis B'? We have that $w(B') \geq w(B)$; we can get B' from B by exchanging the elements that differ, and the one in B' is at least as heavy as that in B (why?!). We cannot have $w(B') > w(B)$, since B is a max-weight basis.

So we will have obtained a third max-weight basis that agrees on a longer prefix with A than B did. We can do this again, getting a max-weight basis B'' differing from B' by a single edge swap; eventually, we will end up with something equal to A. We will therefore have exhibited a sequence of the desired form: Something starting at B and ending up at A, where each sequence is a max-weight basis differing from the previous by a single element.

Corollary to above: all max-weight bases have the same distribution of weights (that is, taking just the element weights and sorting it will yield the same list). (You might think that, e.g., you can have two minimal spanning trees such that one has an edge of weight 10 and 5, and the other has instead two edges of weight 8 and 7; this corollary rules out that possibility.)

Another corollary: Every max-weight basis is a possible output of the matroid greedy algorithm (that is, will be produced if ties in weight are resolved appropriately). You might think that, in Kruskal's algorithm, for example, that you can have a graph with a couple MSTs such that one of them will never be produced; this isn't possible, by this corollary. There will be a tie-breaking that will cause every MST to be produced. It's pretty easy to see what tie-breaking is appropriate: break ties by giving priority to elements in our desired max-weight basis.

A third corollary: if the weights are all distinct, then there is a unique max-weight basis. Cool!

6.2   Matric Matroids

There is another class of matroids called **"matric matroids"**, defined in terms of a given matrix A. Define such a matroid $M = (S, \mathcal{I})$, with S the columns of A, and $\mathcal{I}$ the sets of linearly independent columns of A. You can't model all matroids as matric matroids, but a bunch of the classes of matroids discussed above can be modeled as matric matroids.

It's useful, in general, to consider these special classes of matroids, rather than considering matroids in full generality, because we can often get faster algorithms for specific types of matroids.

7   2/7/2013 (cont): Amortized Analysis

Amortized analysis is typically used in settings where the cost of operations is nonuniform: we ask, if we spend k operations, what sort of bounds can we put on performance. So we may have some operations that are much more expensive than others, but nevertheless we may be able to bound the overall cost as being less than the pessimistic estimate (wherein we upper-bound every single operation by the worst-case performance).

There are different methods of amortized analysis:

1. **Aggregate method**: Simply sum up the cost of all the individual operations to bound the total cost of a sequence of operations. This is the most basic approach, but can get kind of messy.

2. **Accounting method**: Suppose we're trying to prove a theorem that any sequence of k operations will take $O(k)$ time. Then, when an operation is performed, we'll give an operation a constant amount of dollars in order to do it. For example, we'll give an operation \$10 to fund its execution, if we're trying to bound the operations at amortized \$10. The actual semantics are that an operation uses the money we give it; if it uses less than we give it, it puts the rest of the money into different piles of money in the data structure. When an expensive operation takes a lot, then it uses the leftover piles of money from the preceding operations.

   If the money we inject into a system is enough to pay for all the operations, then we can get bounds on the running time.

   Note that it'll frequently make sense to convert this approach to the potential approach, described below (where the potential corresponds to "amount of money left in data structure").

3. **Potential method**: (sometimes, potential function method): A more general method, emphasized more in this course. We define a potential function mapping the state of the data structure to a (usually nonnegative) number. The amortized cost of an operation will be defined as the actual cost plus the change in potential.

   So if an operation changes the state of the data structure from D to D', then the change in potential is $\phi(D') - \phi(D)$, with ϕ the potential.

   We'll use the potential function to make sure the amortized costs all look similar, making it easy to add up the amortized cost.

   This seems weird—what we really care about is the total cost, rather than the total amortized cost (which is this funny function of the actual cost that we made up). However, summing things up makes it make sense:

   Suppose the state of our data structure before/after each of our operations is $D_0, \ldots, D_k$. When we sum the total amortized cost, we get the potential telescoping: the sum of the potentials is

   $$\phi(D_1) - \phi(D_0) + \phi(D_2) - \phi(D_1) + \ldots = \phi(D_k) - \phi(D_0)$$

   which gives us that the total amortized cost is the total cost plus the final potential ($\phi(D_k)$) minus the initial potential ($\phi(D_0)$).

   Usually, the initial potential is 0, and the final potential is typically non-negative; this means that the total amortized cost, according to this method, will actually be an overestimate of the total cost (since total amortized cost = total cost + final potential). So if we can place an upper bound on the total amortized cost, that will also give us an upper bound on the total cost.

7.1 Three Examples

7.1.1 Push/multipop stack

We have a cheap operation, push (push a single item from the stack), and a potentially expensive one, multipop, which pops k times (taking time proportional to k). Multipop(k) can only be performed on a stack with at least k elements.

So a push has actual cost 1, and multipop(k) will have cost k. What is the cost of a sequence of n operations starting with the empty stack?

- **Naive approach**: At all times, we cannot have more then n elements on the stack. No single operation costs more than n: at most, we can pop n elems at any time. Since we have n operations, the total cost is $\leq n^2$.

  This is a very weak upper bound: we'll see that our handful of expensive operations must be offset by a bunch of cheap operations.

- **Aggregate Method**: The total cost of all multipops must be $\leq$ the total cost of all pushes. If we ever multipop off 10 items, we can only do that if, previously, we'd pushed all of them on, at a total cost of n. Since the total cost of all pushes is $\leq n$ (we perform at most n pushes of cost 1), we have that the total cost is at most $2n$.

- **Accounting method**: Give the system \$2 for each push, and \$0 for each multipop. We'll want to show that the pushes have amortized cost 2 and the multipops have 0 amortized cost; this would imply that the total amortized cost is at most $2n$.

  When a push happens, it has to take one dollar; it puts the extra dollar onto the pile. At any point, a stack of m elements will have \$$m$ sitting around. In fact, we can think of each element keeping its dollar, and so we can easily prove that each element has one dollar associated with it, and so we'll be able to fund every operation.

- **Potential method**: Define ϕ of a stack to be the number of elements on the stack (notice that this is basically the same as the accounting argument). Then it's easy to show, using the argument we just gave, that we can bound the cost of operations linearly.

7.1.2 Incrementing a binary counter

We have a counter that has the number 0 initially, and we have only one operation: increment counter. We increment in the most elementary way: start at the LSB, flip and possibly carry, and continue until you hit a 0. So the increment operations are occasionally expensive (if you have to flip a lot of 1's to 0's), but they're usually cheap. We'd hope that, over a large sequence of k increments, our total cost would be $O(k)$.

Concretely, we imagine our counter having an infinite number of bits going off to the left. The first increment will take cost 1; the next will take 2 (we have to flip 2 bits); the third will flip 1 bit; the next 3; and so on.

- **Aggregate method**: the total cost of n operations will give us taht roughly $1/2$ the operations will have cost 1 (every other number is even); $1/4$ of the ops will cost 2; $1/8$ will cost 3, and so on. So the total cost is

$$(n/2) + 2(n/4) + 3(n/8) + \ldots \leq \sum_{i \geq 0} \frac{i}{2^i} = 2$$

(the last equivalence is somewhat nontrivial)

- **Accounting method**: \$2 / operation. We'll need to maintain a different invariant from the stack: we might think that we'd have a dollar per bit, but that won't end up working. We instead have \$1 for each 1-bit in the counter: if our number is 001011, then we have \$3. It's not tough to see this works: incrementing the counter, if we have k bit flips, we had to convert $k-1$ 1's to 0's; we grab that dollar going around, and then deposit our extra dollar on the 1.

- **Potential method**: ϕ will is the number of 1s present in the counter. The amortized cost is the actual cost $+\Delta\phi$. So note that the counter always looks like 01^k for some k (with some stuff before it); an increment flips that to 10^k. So the change in potential is $1 - k$: The potential before is however many 1's there are; the potential after is $1 - k$ minus that (we flip k bits and change 1 to a 0). So the actual cost is $k + 1$; the change in potential is $1 - k$; thus the actual cost $+ \Delta\phi$ is 2. We can thus bound the cost of n increments by $2n$.

  Note this is basically also an accounting argument.

7.1.3 Dynamic Array

The canonical example: suppose we're using an STL vector or a Java ArrayList. The underlying implementation doesn't know how much space is needed; we'll reserve space, but from time to time, the application will exceed the bounds, and a large cost will be incurred. So sometimes a single push can be very expensive, but we'll want to determine that the total cost of a sequence of operations is, say, $O(k)$.

We want space usage to be proportional to the number of elements in the array; that is, we don't want to just allocate way more space than we'll ever need and call it a day. Further, we'll want to ensure that we're using at least a constant fraction of our currently allocated array. Further, we'll want contiguous elements in memory—we can't just allocate space on the heap for each element individually.

We'll do the usual thing—when we run out of space, we'll call `grow`, which doubles the size of the allocation and copies the old elements over. We'd like to give a potential function argument to prove that the amortized cost of grow is $O(1)$. That is, we want a function such that amortized cost = actual cost $+\Delta\phi$ is constant, no matter what phase we're in.

So supposing we have k things in our array, and we have to call grow, we'll have $\Theta(k)$ cost (say, ck). We then need a potential function such that $\Delta\phi$ is $-ck + O(1)$. We'll make ϕ depend on the current size s (the number of elements in the array) and capacity n (note $s \leq n$).

If we let
$$\phi = c(2s - n)$$
we note that during a grow operation, we get $\Delta\phi = 2c$, or somethign? Maybe I got that wrong (go through this argument!) Note that after a grow operation, $s = k + 1$ and $n = 2k$. Note also that ϕ is nonnegative except possibly for the very beginning (which we could handle with a special case)

Suppose now we want to add a `shrink` operation: when the array gets less than a quarter full (that is, $S < n/4$), we want to halve its capacity. Note we cannot just say "halve the length of the array when it's less than half full", because this would give us a sequence of extremely expensive operations: we could go back and forth inserting and removing, right around the point where an array is full, doubling and halving the array, giving us a large fraction of operations that are very expensive.

How can we modify ϕ such that, with shrink, we still get good performance? We design a function such that, when we grow, our potential ramps up from 0 to $\Theta(n)$ (as the previous one did). However, we'll also require that when you shrink from half-full to quarter-full, we want the potential to increase to $\Theta(n)$ too (which our old one didn't do). COnsider

$$\phi = \max\{c(2s - n), c'(n - 2s)\}.$$

The first covers the case we had before; the latter covers the case where we're less than half full (the c constants are different because grow and shrink could have different costs). This function works—when we grow, we get a big change in potential; when we shrink, we get a big change in potential; in between, when incrementing or decrementing, the potential doesn't change much—it's just $O(1)$. The overall amortized cost, therefore, of these "noncritical" operations is constant. (Note that sometimes the amortized cost will end up 0 or negative; that's fine, just a side-product of the funny math going on.)

8   2/19/2013: Splay Trees

8.1   Motivation, definition

Recall an ordinary Binary Search Tree. There, we have the invariant that, at any subtree, the value of its root is $\geq$ any value in its LHS subtree and $<$ any element in its RHS subtree. Doing an inorder traversal gives us the elements in sorted order.

The operations that we want to be able to insert are Find, Insert, and Delete. In general, these can be expensive: in the degenerate case, a BST can devolve into a linked list, and Find/Insert/Delete can take linear time. That is, these operations are all $\Omega(n)$ worst case.

We therefore consider various "balanced" binary search trees, like red-black trees or AVL trees. They're similar to BSTs, but various "rotation" operations are used to maintain $O(\log n)$ depth. A rotation is basically like picking up a tree by a node and rebalancing: swap a child and parent (and possibly move their subtrees), for example. These give find/insert/delete performance of $O(\log n)$ in the worst case.

Splay trees give $O(\log n)$ amortized time for find/insert/delete. They don't give guarantees about individual operations like the balanced trees, but we get good total performance (that is, the total cost gives the same upper bound as what you'd get with a red/black tree, but individual operations may be expensive).

Splay trees will end up being fun because the working set gets moved to near the root of the tree, so we'll get something that's actually less than the log of the number of nodes, but instead log of the working set runtime (which will almost certainly be smaller).

What is a splay tree?

- **Find**: First, do an ordinary BST find, and then do a splay operation. A **splay operations** is a sequence of **splay steps**. There are three different kinds of splay steps:

  1. A standard rotation operation. Suppose we have y the root. Suppose y has LHS child x and RHS subtree γ. Suppose x has LHS/RHS children α, β.

     We bring x up to the root position, making α its LHS subtree, y its RHS child. We make y's LHS/RHS children β, γ. Note y must be the root.

     The symmetric case does the same.

  2. A "zig-zag" step. Suppose we have some z, and x is the right child of the left child of z. Call the left child of z y.

     What we do is put x at the top, putting y as its LHS child, z as its RHS child. We move their subchildren appropriately (x's left subtree becomes y's right; x's ight subtree becomes z's left.

  3. "Zig-zig". Supposeing we have x the left child of z's left child (which is y). Then we move x to the top, putting y on its right, and z on y's right (and moving the subtrees as appropriate).

  A splay operation at x, then repeatedly performs splay steps on x until it is at the root. If x is a child of the root, then we do the first thing. If x is at least 2 levels away from the root, then we will land in exactly one of the four above cases (the two symmetric cases of each of zig-zig and zig-zag).

  Note that every one of the splay steps above retained the tree's BST property.

  If we have a successful find in the BST-find, we perform splay operations to move the node to the root. If we don't (that is, x isn't in the tree), then it will have failed at some node y; we splay y up to the root.

Why do we splay y up in the case of a failure? Well, if there's a long path, then, intuitively, we don't want to keep the long path around. By splaying y to the root, we'll eventually make the tree more shallow.

- **Insert**: follow downward from the root, like you're doing a find. Once you reach a point where you would branch into an empty subtree, just insert it there. We then splay the newly inserted item to the root. That is, we again do the BST thing up to the root.

- **Delete**: This one's a bit more hairy so we don't talk about it here. Suffice it to say, we follow the same basic procedure—do a classical BST operations, and then splay something up to the root (the parent of the deleted thing, say).

8.2 Amortized analysis of find

We give a potential function argument to argue that a series of find operations is amortized logarithmic. Defined the rank of x, denoted $r(x)$, as the floor of the base-2 log of the size of the subtree rooted at that node. Note that "size" here means "number of nodes", not "height".

Define the potential Φ as $c \sum_x r(x)$, summing over all nodes x (we determine what we need c to be later).

We think of the amortized cost of splay-tree find as the amortized cost of phase A + the amortized cost of phase B (the former being the BST find, the latter being the splay operation).

Assume that find reaches depth d (with $0 \leq d < n$). Then the amortized cost of phase A is the actual cost of phase 1 plus its change in potential, which is $c_0 d + \Delta\Phi$ in phase 1; since phase 1 didn't change the potential, $\Delta\Phi$ is 0, and our amortized cost is $c_0 d$ (for some constant c_0 bounding the cost that we spend in each level of the tree).

The amortized cost of phase B is $c_1 d + \Delta\Phi$ in phase 2. $c_1 d$ is the actual cost of the about-$d/2$ splay steps we need to perform. We must show that $\Delta\Phi$ in phase $2 \leq (c_2 \log n) - (c_0 + c_1)d$. That is, when we add up the amortized costs of A and B, we need to get at most some constant of $\log n$, so we need to wipe out the $c_0 d$ and $c_1 d$ terms present in the actual costs of the two steps.

We examine the sum of potential changes due to the individual splay steps performed in phase 2. For a splay step at step x, define $r'(x)$ as the rank after the step, and $r(x)$ as the rank before. We'll focus on the two zig-zag steps; we perform at most 1 of the steps that occur at the root of the tree.

How can we get at the change in potential of a zig-zag operation? Consider the third operation, zig-zig. Note that the potential doesn't change at almost every node in the tree. The only nodes where the contributions to the potential may have changed are the nodes x, y, and z involved in the zig-zag. The chagne in Φ is

$$c(r'(x) + r'(y) + r'(z) - (r(x) + r(y) + r(z)).)$$

Note that $r'(x)$ and $r(z)$ cancel out—they're both sitting at the root of the tree. (The same reasoning holds for zig-zag.)

We partition zig-zig and zig-zag steps into "type A" steps and "type B" steps, and then look at what their sums of potential changes looks like. Type A steps are those where $r'(x) > r(x)$. **Claim**: the total change in potential due to type A splay steps is at most $c_3 \log_2 n$ (for some c_3).

Consider again the change in potential:

$$c(r'(y) + r'(z) - (r(x) + r(y))).$$

We have that $r'(y)$ and $r'(z)$ are each at most $r'(x)$ (x is at the top of the subtree after the operation); we similarly have $r(y)$ and $r(x)$ being at most $r(z)$. So our change in potential is less than

$$\leq 2c(r'(x) - r(x)).$$

That is, the change in potential is upper-bounded by twice the increase in rank.

Note the $r(x)$ is initially nonnegative, and it can never go below $\log_2(n)$. In fact, the final rank of x is (upper-bounded by)

$$\lfloor \log_2 n \rfloor.$$

Summing over all the type A steps (the ones that increase rank), what is an upper bound on the sum? Well, it is at most $2 \log_2 n$ (that is, we've proven the claim with $c_3 = 2$). (I am missing an important statement in here somewhere.)

Intuitively, the case we're worried about is where we have a linear number of splay steps. What we've shown is that we can have at most a logarithmic number of these, since each of them causes an increase in rank.

The vast majority of splay steps will be type B, where $r'(x) = r(x)$. This happens because we take the floor of $\log_2$; unless we get over a power of 2 (which doesn't happen much), the rank doesn't change.

Claim: each type B splay step gives a potential change $\leq -c$ (for some c). This is really the heart of the argument: we're relying on the type B steps to give us a large drop in potential.

Recall the claim above that the change in potential is $\leq 2c(r'(x) - r(x))$. This isn't strong enough; in the type B case, this basically shows there is no increase in potential. What we really need to argue is that

$$r'(x) + r'(z) - r(x) - r(y) \leq -1.$$

Consider zig-zag. Suppose we have $z- > y- > x$; x ihas subtrees β, γ; y has LHS subtree α; z has RHS subtree δ. Let s be the total number of nodes in this subtree. Let t be the number of nodes in the tree rooted at x (that is, x with β and γ). We have that $t > s/2$; if it were otherwise, then the rank at x would have gone up in the element.

Now, consider the thing after the zig-zag, rooted at x. Either a or b is $< s/2$ (otherwise there would be more than s elements in the tree rooted at x).

27

Suppose that $a < s/2$. Note either $r'(y) < r'(x)$ or $r'(z) < r'(x)$ (why?). That gives us that $r'(y) + r'(z) < 2r'(x)$ (why?). This gives us our desired result (why?).

A similar line of reasoning can be applied to zig-zig. We have $z-> y-> x$, and we switch it to $x-> y-> z$, with initially x the lhs of the lhs of z, and eventually z the rhs of the rhs of x.

Let the whole thing contain s nodes; let t be the size of x with α, β its subtrees. Now, we have the RHS-most subtree after the step (z with the otehr subtrees) have strictly less than half of s. this ends up giving us that $r'(z) < r'(x)$ (recall x ends up at the root of the resultant subtree, with LHS subtree α and rhs node y, with LHS subtree β and rhs child z.

(end proof of claim)

We DO have to say something about the last step. It's pretty simple: basically, we show we don't get a massive explosion in potential with the operation; at worst, we get at most a logarithmic term.

We have at most $\log_2 n$ Type A splay steps, out of about $d/2$ splay steps (since each gives us a bump up in rank of about 1). Thus, the number of type B splay steps has to be $\geq d/2 - \log_2 n$. Thus, the total type B drop is at least $cd/2 - c\log_2 n$; this is the claim we wanted to show above. That is, we can show that the drop in potential is $\leq c_2 \log n - (c_0 + c_1)d$. We get this by setting c_2 to $3c$, and choosing c such that $c/2 = c_0 + c_1$.

————————————

For insertion and deletion , the argument ends up being very similar. Note inerstion is tricky: we potentially add anotehr element, and so every single node on its path might have its potential increased. A naive argument will give us an insertion perhaps increasing our potential by $\Theta(n)$; however, if we notice that adding a new node can only strictly increase the rank of a logarithmic number of nodes; you wander up to the root, and by the time you reach the top, you arrive at potential at most $\log n$ (why???).

9 2/21/2013: Augmented Binary Search Trees

9.1 Example 1: Dynamic Order Statistics

Simple technique, fairly widely applicable.

If you have a dynamic set of items drawn from some ordered universe, we can use Red-Black trees or AVL trees or something like that to get logarithmic time for insert/delete/find. However, sometimes we want one or more additional operations, typically more specialized, and we want logarithmic performance. Instead of just starting from scratch, we can frequently just augment a standard balanced search tree (e.g. red/black tree).

As a simple example, consider the problem of "dynamic order statistics". Suppose we have a dynamic collection of elements, but from time to time we'll want to run SELECT(k) queries, which returns the kth highest element in the set. If we just use a red-black tree, we can't do this efficiently—we'd have to

do an in-order traversal to figure out where in the tree we want to be, and that would take linear time. However, we would like to get logarithmic time for this.

We can do so by augmenting a red/black tree in such a way (adding extra fields) that will be helpful for this. We maintain a size field at each node, which is kept equal to the number of nodes in the corresponding subtree. We'll need to work out exactly how to easily maintain this field, but it ends up being possible.

We'll have some size k_0 at the root. Say k (the query value) goes from 1 up to n. Call k_1 the size field of the root's left child (if there is no left child, then the value is 0). If the rank we want is $k_1 + 1$, then we return the value at that node. If we want a rank that's less than that, we recurse into the left tree's subtree. If we want a query of rank greater than $k_1 + 1$, we recursively search in the root's right subtree, with $k' = k - (k_1 + 1)$.

Note we never descend into both sides of any subtree; since a red-black tree has depth $O(\log n)$, we will spend logarithmic time on the query.

9.2   Augmentation Theorem

We have the following **augmentation theorem**, which is proved in CLRS (chpt 14). Suppose the auxiliary fields of a node x can be computed in constant time given the elements in x and the auxiliary+element fields of its children (that is, we can compute the auxiliary info for a node, given all the info present in its children, and the original, non-auxiliary info at that node). Then we can maintain these auxiliary fields, while achieving logarithmic performance on a red-black tree for insert/delete/find.

We give a sketch of a proof of this. We need to show that we can augment the code for insert/delete/find in such a way that we maintain logarithmic time. For find, we don't modify any auxiliary fields, so the performance is trivially the same. Now, insertion. Recall that red/black tree insertion just performs a typical BST insert (putting the new value as a leaf); in the second phase, we perform a constant number of rotations. In the first phase, we may have to percolate information up to the root, but since the tree is balanced, that will take logarithmic time. In the second phase, we may have to propagate a change up to the root again. Again, this is at worst $O(\log n)$.

This argument doesn't apply to something like AVL trees, where you might do a logarithmic number of rotations on insertion. However, we can amend the argument—you won't be doing rotations anywhere in the tree; all of them will be along a path to the root, and so you can prove that there is still only a logarithmic amount of extra work done.

The argument for deletion on red-black trees is very similar.

9.3   Example 2: Finding an overlapping interval

Suppose elements in our dynamic sets are intervals (a, b) where $a \leq b$. We want to support an operation OVERLAP(a_0, b_0) (representing the inclusive interval $[a_0, b_0]$), that either finds an overlapping interval in our dynamic set, or reports correctly that there is no such interval. How do we augment the tree?

We sort the intervals by left endpoint in the red-black tree. For example, suppose we have a root with two kids. Root has a_1, b_1; the LHS kid has a_0, b_0, and the RHS child has a_2, b_2. Suppose we have the query (a, b); what do we need to augment the tree with?

Suppose that $\text{OVERLAP}(a, b)$ happens to overlap the root. Then we're done. If it doesn't, however, we want to query at most one of the two children. If a is between a_1 and b_1, we're done. We're left, then, with the two cases of $a < a_1$ and $a > b_1$.

Suppose $a > b_1$. If the maximal right endpoint of any interval in the LHS subtree is greater than a, then we have an overlap; if it's less than a, then every interval in the LHS subtree is guaranteed to fall completely to the left of a, and therefore there will not be any overlap.

We therefore say that our auxiliary field can be the max b-value in a subtree. If we have that, we can compare a to the max b-value of the LHS tree (assuming $a > b_1$); we can therefore recurse to the right, since we know that there is no possibility of finding an overlapping interval to the left.

Suppose now that $a < a_1$. Then we want to know what b looks like. If $b > a_1$, then we've found an overlapping interval. If, on the other hand $b < a_1$, we know that it's not possible that we get any intersection with any of the intervals in the right subtree (since the tree is sorted by a value). We can therefore recurse left.

9.4 Example 3: Finding (x, y) with $x > x_0$ and largest y

Say our elements are arbitrary (x, y) pairs in an ordered universe. Given a query x_0, we want to compute an element (x, y) with maximum y-value over all elements with x-value $\geq x_0$. For example, the x's could be release times, and the y could be values; we'd then be looking for the highest-value object after a certain release time.

Looking at it, it's not obvious if we want to sort by x or by y. It turns out you can get pretty nice solutions either way.

Solution 1: Sort by y-value. Say we have a tree with root x_1, y_1, LHS child x_0, y_0, and RHS child x_2, y_2. We're given a query with argument x.

We want to branch to the right whenever possible (we're trying to maximize y). The only time we don't want to do that is when we can prove there are no useful points to the right. When does this happen? Exactly when all of the x values in the RHS subtree are too small. We can determine this efficiently if we keep around the maximum x-value of a subtree as an auxiliary component.

When we query on x, we look at the RHS node's max x value. If it's not big enough, then we know there is no point in going that way. We can then check if the node's value itself is acceptable (in which case we return the root). Otherwise, we branch left.

Solution 2: Sort by x-value. We need to perform two passes. The first pass determines the y value; the second determines (x, y). The auxiliary field will be the max y-value.

What we want to do to find the max y value is the following. Supposing we're at the root, with a query x, we check to see if $x > x_1$. If so, then that node's acceptable, and all RHS tree nodes are acceptable. We can branch left, then, but remember the highest y value in the root or in the RHS subtree (we do that by taking the max of the auxiliary fields in the root and the RHS node). We then recurse to the LHS subtree. If $x_1 < x$, then we know that we must branch to the left—neither the root nor the LHS subtree contain any feasible nodes.

We thus come to know the y value we're seeking, and it's not too tough to figure out where to find the x value (what? how?)

9.5 Example 4: Sum all y values from (x, y) pairs with $a \le x \le b$

Suppose we have pairs (x, y) at nodes, and the query is (a, b). We want to determine the sum of the y-values of all elements (x, y) such that $a \le x \le b$. We sort the (x, y) pairs on x, breaking ties arbitrarily. Our auxiliary value is the sum of y-values in the subtree.

We give a two-pass solution. First, we traverse the find-a path (that is, run a BST find on a), computing the sum of all the y-values of all the points with x-value at least a (we can add up subtree sum values). Second, we find the sum of the y-values of all points $\ge b$. We can just subtract the latter from the former, then.

9.6 CORRECTION

In order to be able to support find efficiently, in the above, we're actually going to need to sort on lexicographic order on the (x, y) pairs (that is, break ties on the second field). Imagine the pathological case of a tree with all x values the same, for example.

10 Van Emde Boas Trees

(This is Chapter 20 in the 3rd edition of CLRS.)

A Van Emde Boas tree is a data structure that can substantially outperform red-black trees in some situation. Following the text, we want to support an operation MEMBER (return 1 if key in set, 0 otherwise); we won't be considering auxiliary data, but we could perhaps also store auxiliary data, and modify FIND to return a pointer to it. That is, we're building something that does set membership, rather than an efficient map. We also have MINIMUM; MAXIMUM; SUCCESSOR; PREDECESSOR; INSERT; DELETE.

We'll make a set of restrictive assumptions about our keys so as to be able to get sublogarithmic time on our operations. In particular, we assume the integers are drawn from a set $\{0, \ldots, u - 1\}$. We'll eventually get $O(\log \log u)$ time for all operations (in fact, min and max will be $O(1)$).

First, note the simple naive bit-array version: Represent a set as a list of u bits. Then MEMBER, INSERT, DELETE are all $O(1)$. However, finding the SUCCESSOR is $O(u)$ (we could have to examine many bits before finding a 1). Similarly with PREDESSOR, MIN, and MAX.

A variation of the above would be to superimpose a binary stree structure on the bit vector: We could store a $\log u$ binary tree, where an intermediate node of value 0 means "there are no leaf nodes with 1 below this node". This gives us $O(\log u)$ SUCCESSOR, PREDECESSOR, MIN, and MAX, but the other operations won't be $O(1)$ any more. However, this doesn't really give us much over a red-black tree, which already gave us $\log n$ time, rather than $\log u$ time.

Another variation could be to use a k-ary tree, for example a $\sqrt{u}$-ary tree. This would give is OK time (figure out what it would be).

We'll use a branching factor proportional to the size of a subtree. So at the top, there are u leaves; we create $\sqrt{u}$ child nodes at the root. Each child has $\sqrt{u}$ nodes; we give each of them $u^{1/4}$ children; since there are $u^{1/2}$ of them, there will be $u^{3/4}$ nodes at that level. The depth of this tree is $O(\log \log u)$. We get this because at each level there are $u^{1/(2^i)}$ nodes (FLESH OUT THIS A BIT MORE).

10.1 Proto Van Emde Boas trees

Following CLRS, for pedagogical purposes, we first introduce a version of van Emde Boas trees that won't work out for us.

A size u node will have $\sqrt{u}$ children, each of them built in the same way (as above). However, we also have a summary of size $\sqrt{u}$ at the node—since we already have children of size $\sqrt{u}$, this won't change the size of the trees in big-O terms. The summary will look like the following. We are taking care of a subset $0, \ldots, u-1$. We imagine chunking this up into blocks of size $\sqrt{u}$. The first such one will be bits $0, \ldots, \sqrt{u} - 1$ (Call this cluster 0). Cluster 1 will be the next $\sqrt{u}$ bits. We continue this way up to cluster number $\sqrt{u} - 1$. Our summary will have a bit for each cluster, which is just the bitwise-or of all the bits in a cluster (that is, the semantics are "is at least one bit in the cluster a 1?").

We use the following notation from the text: when we're at a node with size u, the value x in the query (for example, MEMBER), will be a $\log u$-bit value. We use $HIGH(x)$ to refer to the most significant $(\log u)/2$ bits, and $LOW(x)$ to refer to the least significant $(\log u)/2$ bits. The HIGH bits will tell us what cluster to go into, and the LOW bits will tell us the element in the set determined by the HIGH bits.

At a node, doing an operation, we'll use the summary to skip clusters that we know have no bits in them. Supposing we're using $u = 2^{16}$, at the root node we'll have 2^8 clusters, and the summary is a representation of the integers $0, \ldots, 255$; if i is in this set, then the ith cluster has at least one number in it. Note that each of the honest-to-god children is also representing a subset of the integers 0 through 255 (that is, 0 through 2^8, representing the range of LOW bits). So the summary child looks just like the other children.

At the next level, clearly we'll look the same: each node has 2^4 children nodes and summary of 2^4 bits.

How does MIN work? The natural way is to recursively ask the summary what its MIN is; supposing that the summary's MIN is 17, then we know that we can ignore clusters 0 through 16, and we can recursively ask cluster 17 its MIN. WE have made 2 recursive calls on data structures of size $\sqrt{u}$ to get the answer of a structure of size u. This gives us the following recurrence for MIN with proto-VEB trees:

$$T(u) \leq 2T(\sqrt{u}) + O(1). \tag{20}$$

We also have the notation $INDEX(HIGH(x), LOW(x))$; when we recursively get an answer from a child, we need to concatenate two half-words together. So if we recursively get an answer at cluster k, and k returns ℓ, then our answer will actually be the concatenation of the high bits k and the lower bits ℓ. The book denotes this as $INDEX(HIGH(x), LOW(x))$ or something.

Unfortulately, unfolding recurrence (20) gives us

$$O(2^{\log \log u}) = O(\log u)$$

as its runtime (if we'd had $T(n) \leq T(\sqrt{u}) + O(1)$ instead, THEN we'd get $O(\log \log u)$). So we'll have to change our data structure to get us the desired behavior (in particular, we'll want to change the 2 coefficient, corersponding to the 2 MIN calls made, into a 1).

Similarly, what does SUCCESSOR(x) look like? First, we look in HIGH(x). If we found a successor there, we're in luck. However, if HIGH(x) is empty, then we'll have to look at the next cluster to the right that has anything in it. So we'll need to make a recursive call to the summary node, and find the MIN element in that (which we already looked at). So we'd get

$$T(n) \leq 2T(\sqrt{u}) + O(\log u) \tag{21}$$

the $O(\log u)$ is from the MIN call. This ends up being $O(\log u \log \log u)$, which is again not really what we want.

Consider now INSERT on the proto-VEB trees. What do we do? We insert LOW(x) into cluster $HIGH(x)$. If we're unlucky, the cluster HIGH(x) was empty before, and we'll need to insert into the summary node the fact that the cluster is now nonempty. That gives us

$$T(n) \leq 2T(\sqrt{u}) + O(1)$$

which is $O(\log u)$.

Note that I think MEMBER is still $O(\log \log u)$ here.

10.2   Van Emde Boas Trees: fixing things up

10.2.1   A Modification

What we need to do is patch things so that we don't need the second recursive call. What we do, intuitively, is change things so that we have two cases: an

easy case and a hard case. We'll have good behavior in the easy case, and kind bad behavior in the hard case; when we need two operations, then, we'll rig things so that only one of them is "hard".

We'll maintain a MIN and MAX at each node (setting them to Nil if there are no elements). Note that, at a node, the Minimum element is not represented below.

10.2.2   An Example

As an example, suppose our set is $\{2, 3, 4, 5, 7, 14, 15\}$; then $u = 2^{2^2} = 16$. At the root, MIN=2 and MAX=15; we won't keep the 2 in its children.

The root has 4 children and the summary. The 4 clusters are $0-3, 4-7, 8-11, 12-15$ (clusters 0 through 3). Node 0 has $u = 4$; it has to represent the elements from $0-3$, except for the min (and is therefore really a representation of the set $\{3\}$. Cluster 1 is suppoesd to keep track of the ints between 4 and 7. Since it only cares about the lower-order bits, 4 gets mapped to 0; 5 to 1, and so on. So it has the set $\{4, 5, 7\}$, which is represented as $\{0, 1, 3\}$. The parent node, when it gets an answer from cluster 1, will add 4 to the result, since it knows it's cluster 1. And so on.

The summary has $\{0, 1, 3\}$, since clusters 0,1, and 3 are nonempty, while cluster 2 is empty.

We have similar tree structures for each of the above-mentioned subtrees. Each child has 2 clusters (of size 2) and 1 summary; each cluster has size 2, representing a subset $\{0, 1\}$. Recall the MIN of any node is not stored in its subtree.

10.2.3   Performance

Trivially MIN and MAX are $O(1)$ (they're both stored directly at the root). MEMBER does the following: if the value you're looking for is equal to MIN or MAX, then you know it's in the set. Otherwise, split x into HIGH(x) and LOW(x) in constant time; we ask cluster HIGH(x) if LOW(x) is a member of it; that'll give us $O(\log \log u)$ time.

A bit tricker is SUCCESSOR. We have the trivial cases with $u = 2$ (resolvable in constant time), or SUCC = MIN, in which case we can immediately return the MIN (Note that the successor is defined as the smallest value in the set that is greater than the query integer x. FLESH THIS OUT). We want to show that we can now get away with a single general recursive call.

What we do first is find MAX-LOW: find the maximimum value in cluster HIGH(x). If MAX-LOW $\neq$ NIL and $LOW(x) < MAX - LOW$, then we know that our answer must be within cluster HIGH(x), and we can just find the successor of LOW(x) inside the cluster HIGH(x), then build the answer (which is HIGH(x) concatenated with the answer we got before). If, on the other hand, either there's nothing in that cluster or the answer is in a higher cluster, then we have the same problem we had before: we look for the successor of HIGH(x) in summary (that is, find the next nonempty cluster to the right). This is $T(\sqrt{u})$.

Then we get the answer as the MIN value in the successor cluster (formally, that concatenated with the successor).

WHY DON'T WE KEEP THE MIN?

For the code of INSERT and DELETE, consult chpt 20 of CLRS.

11   2/28/2013: Union Find

Here, we'll get something growing much more slowly than even $O(\log \log u)$; in fact, we'll get an inverse Ackermann-growing function.

The problem of Union-find is the following. Suppose we want to maintain a dynamic collection of disjoint sets. Initially, we have no sets. WE have an operation called MAKESET(x), taking an elemennt from some universe, and creates the singeton set $\{x\}$ (we require that the element x not already be in some element in our collection—we require disjoint sets). Further, we have an operation UNION(x, y), which merges the sets containing x, y; that is, there should be a set in our collection with x, and one with y; after the call, we'll remove the set containing x, remove the set containing y, and then insert the union of those two sets. Finally, we have FIND(x), which returns the "name" of the set containing x (we require that x be an element in some set in the data structure already). It doesn't really matter what this "name" is; we only require that x, y are in the same set iff FIND(x) = FIND(y).

Tarjan discovered the following union-find data structure, which uses trees to represent each of the sets. Each element of a set is a node in a tree; each element has a pointer pointing to its parent (there are only upward-pointing arrows in the tree). For any node, we simply follow the upwards parent-pointers until we get to the root. Note there are many possible representations of a set; all we require is that it be a tree. The naming convention will be to return the root element as the name of a set.

Note that on a find call, we're given directly the node in the tree; we don't have to look around in different trees to find an element.

If we want to do a UNION on two elements in the same set, we don't need to do anything. If they're in different trees, we do the following: make one of the root nodes a child of the root of the other one. Following CLRS, define LINK(T, T') as an operation taking as arguments two roots of trees; this operation makes one of them point to the other (that is, one of the roots is a child of the other root).

Given what we've defined, we can imagine degenerate cases with bad behavior. In order to get around these, we have two tricks.

First, we have **path compression**. Whenever we do a FIND, we traverse the "find path"—the path from the found node to the root. With path compression, we retraverse the path a second time, and make everything point directly to the root. That is, if we traverse a path of length 100 to get to the root, we make 99 elements that didn't used to point to the root point directly at it. Subsequently, FIND on any of those elements will be very cheap (assuming we haven't done too

many intermediate UNION operations). This gives us pertty good performance by itself.

Second, we have **union by rank**. With this operation, we'll be able to almost get amortized $O(m)$ time of any sequence of m operations. In fact, what we'll get instead is m operations taking $O(m\alpha(m))$ time, with m the number of MAKESET calls, and α the inverse Ackermann function; however, the latter is very very slow.

Anyways, back to **Union by Rank**: if you have a huge structure and a tiny one, it'll be better to make the tiny structure a child of the large one. Intuitively, if we make the large structure a child of the small one, then we would increment the depth of a whole bunch of nodes; if we make the small one a child of the large one, we only increment the depth of a small number of nodes.

Instead of making this decision by size, we'll make it by **rank**, which is a proxy to size, which will end up giving us the same optimal bounds. Rank is a nonnegative integer. When we do MAKESET, initialize the element's rank to 0. A FIND has no impact on rank; when we call LINK(x, y); if x has rank r and y has rank r', and assuming wlog that $r \geq r'$, then we make y a child of x. If $r > r'$, then there is no change in the rank of x; if $r = r'$, then we increment the rank of x. That is, regardless of the structure of the tree, if we link two trees with the same rank, then the rank is the previous rank plus 1; otherwise, the rank is what it was beforehand.

A side note: you can show pretty easily that the $r(x) \leq \log_2 n$, with n the size of the subtree rooted at x; further, the depth of the tree rooted at x is at most the rank of x. This can be shown by induction on the tree size.

Now, the Ackermann function will look as follows:

$$A_0(j) = j + 1 \tag{22}$$
$$A_1(j) = 2j + 1 \tag{23}$$
$$A_2(j) = 2^{j+1}(j + 1) - 1 \tag{24}$$
$$A_3(j) = 2!j \tag{25}$$

where $2!j$ is a tower of repeatedly exponentiated 2's (that's notation people use, right?). I can't type A_4, but it's a tower of 2's, whose height is a tower of 2's, whose height is a tower of 2's, and so on. The functional inverse of doubling is halving, and the functional inverse of exponentiation is log, and the functional inverse of the tower is $\log^\star$, which is something or other (I dunno).

Define $\alpha(n)$ as the minumum k such that $A_k(1) \geq n$. Note the real function used in Tarjan's paper is a bit more complicated (it's the Ackermann function with two arguments).

Formally,

$$A_0(j) = j + 1 \tag{26}$$
$$A_k(j) = A_{k-1}^{(j+1)}(j) \tag{27}$$

where $A_k^{(j)}$ is A_k iterated j times.

We'll see that the running time of m operations is $O(m\alpha(n))$, with n the number of MAKESETs. Note LINK is constant-time worst-case, MAKESET is constant time. The difficulty is in FIND—we can have a long path to the root, but we'll want to show that this doesn't happen too often.

We'll use a cleverly-defined potential function to argue that the amortized cost of each FIND is $O(\alpha(n))$. NOte that LINK will still have amortized $O(\alpha(n))$, even though it's worst-case cosntant time (because it can potentially increase the potential). When we do a FIND, we'll want to argue that the potential function drops in such a way that it basically entirely offsets the cost of traversing the path to the root. Since we pay 1 unit of actual cost going up a step in the path, we'll need our potential go down 1; there will be $\alpha(n)$ nodes in the path in which we can't make that argument, so we'll have to pay the unit cost at each of those nodes.

The overall potential Φ will be

$$\Phi = \sum_x \phi(x)$$

with x ranging over all nodes in the data structure, and $\phi(x)$ being

$$\phi(x) = r(x)(\alpha(x) - LEVEL(x)) - ITER(x)$$

where $LEVEL(x)$ is the following: for a node x, we look at the rank of the parent of x, defined as r'; if r' is much bigger than x, we'll ahve a high LEVEL, otherwise we'll have a low LEVEL (intuitively, it will encode "how much higher a parent's rank is than yours"). Precisely, defining $r'(x)$ as the rank of x's parent, $LEVEL(x)$ is the smallest k such that $A_k(r(x)) \geq r'(x)$. Note that for any non-root x, we have that $r(x) < r(p(x))$ (with $p(x)$ the parent of x) (it is somewhat easy to convince yourself of this). So the question is, roughly, "how much higher" is your parent's rank than yours? Note that LEVEL will almost always be really small.

Below, we abbreviate $LEVEL(x)$ as $\ell(x)$ and $ITER(x)$ as $i(x)$.

LEVEL is a primary measure of how much higher a parent's rank is than yours. ITER is a secondary one. $ITER(x)$ is the largest k such that

$$A_{\ell(x)}^{(k)}(r(x)) \leq r(\phi(x))$$

with $A_{\ell(x)}^{(k)}$ denoting k successive applications of the Ackermann function $A_{\ell(x)}$ (and with $\ell(x)$ the LEVEL function). $\phi(x)$ is

$$r(x)(\alpha(n) - \ell(x)) - i(x)$$

for a nonroot x with $r(x) > 0$; and $r(x)\alpha(n)$ otherwise.

Note some properties of $\ell(x)$: first, $0 \leq \ell(x) < \alpha(n)$. I'm not quite exactly sure why this is, but I think it's basically just an application of the definition of LEVEL. SEcond, the level of a node cannot decrease. We might get a new parent after path compression; however, since the rank of the parent of x cannot

decrease under any operation, and the rank of a parent is always larger than its kids, then the rank of a parent will only stay the same or increase.

Note the following tow properties about $i(x)$. First, while $\ell(x)$ does not change, $i(x)$ cannot decrease either. This follows from a similar line of reasoning for why $\ell(x)$ cannot increase, I think (but I'm not sure). Second, $1 \leq i(x) \leq r(x)$. To show thise, we use contradiction. Suppose $i(x) \geq r(x) + 1$. Note that

$$A_{\ell(x)}^{r(x)+1}(r(x)) \leq A_{\ell(x)}^{i(x)}(r(x)) \leq r(p(x))$$

for some reason; further, the LHS-most term in athe above is equal to

$$A_{\ell(x)+1}(r(x))$$

This is a contradiction for some reason I don't understand at all.

Heading back, note that $\phi(x) \geq 0$ for every x (why?), so our potential Φ is positive; further, $\Phi = 0$ initially.

Let's look now at the amortized cost of a FIND operation. WE could have some sort of really long FIND path (linear in m, say); we'll want to show that the amortized cost is $O(\alpha(n))$. We'll need to offset that cost with a corresponding drop in potential. Why would this happen? Intuitively, for a whole lot of nodes, the comparative rank of the nodes' new parents will be higher than their old parents'.

WE get through some convoluted argument that, for any node x, the drop in potential given by that node is at least 1. I couldn't follow the argument because I am sick and can't focus on anything. It seems like the argument is probably laid out in CLRS though.

At the end of the day, the total amortized cost of a FIND is O(number of nodes x on the FIND path for which $\phi(x)$ is unchanged), that is, for which $\ell(x)$ and $i(x)$ are unchanged; those nodes contribute 0 to the total cost, whereas the other nodes cancel out the unit amount of work needed to visit/move that node.

How on earth can we upper-bound that number? Well, suppose we have a bunch of level 7 nodes, say. Then for every node except possibly the topmost one, we can be sure that the level will drop. Why? well $\phi(x) \geq A_7^{(12)}(r(x))$ but less than $A_7^{(13)}(r(x))$. Then consider every level 7 node except for the topmost one; each of these leapfrogs someone of the same level, and something interesteing happens there. or whatever, jesus i have no idea.

More concretely, suppose we have three rank-7 nodes a, b, c, with a the furthest from the root, b the closest; they're along a path. Now, calling y the new root, we have

$$\begin{aligned}
r(y) &\geq r(c) \\
&\geq A_7(r(b)) \\
&\geq A_7(A_7^{(12)}(r(x))) \\
&= A_7^{(13)}(r(x))
\end{aligned}$$

and I have no idea what this means, but apparently it gives us some insight as to why only the topmost level-7 node might be problematic. This, further, supplies the insight for why the prose-language expression inside the big-O above is $\leq \alpha(n)$.

WE need also think about the amortized cost of a LINK operation. Supposing we link x and y with $r(x) \leq r(y)$, then $\phi(x)$ cannot increase, but $\phi(y)$ could increase, I think by $\alpha(n)$ for some reason. Oh, I think it's entirely due to a possible increase in potential, and has nothing to do with cost. This still gives us what we want, though.

12 Basic Graph Algorithms

12.1 Directed Graph Traversal